



# Statistical Machine Learning

## Lecture 10: Linear Dimensionality Reduction & Statistical Learning Theory

Kristian Kersting

TU Darmstadt

Summer Term 2020

# Today's Objectives

- Make you understand how to reduce dimensionality and how to understand the power of function approximators!
- Covered Topics:
  - Principal Component Analysis
  - Statistical Learning Theory

# Outline

**1. Linear Dimensionality Reduction**

**2. Statistical Learning Theory**

**3. Wrap-Up**

# Outline

**1. Linear Dimensionality Reduction**

2. Statistical Learning Theory

3. Wrap-Up

# Towards Principal Component Analysis

- We have covered a number of linear methods
  - For classification and regression
  - Both are supervised learning settings, i.e., pairs of input/output training data points
- Sometimes, it is quite helpful to analyze the data points themselves
  - Unsupervised learning
  - Particularly: Reduce the dimensionality of the data
  - Possible application: Visualization of the data

# Motivation from Linear Least-squares Regression

- In straightforward least-squares linear regression the parameters are computed as

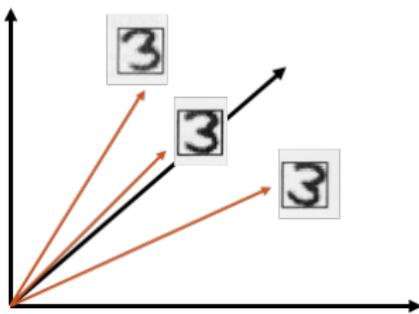
$$\hat{\mathbf{w}} = (\hat{\mathbf{X}}\hat{\mathbf{X}}^\top)^{-1}\hat{\mathbf{X}}\mathbf{y}$$

where  $\hat{\mathbf{X}} \in \mathbb{R}^{d \times n}$  and  $\mathbf{y} \in \mathbb{R}^{n \times 1}$

- We need to invert a  $d \times d$  matrix, which naively costs  $O(d^3)$
- Hence, it would be helpful to find a new  $d^{\text{new}} \ll d$  to gain computational advantage while not loosing prediction performance

# Dimensionality Reduction

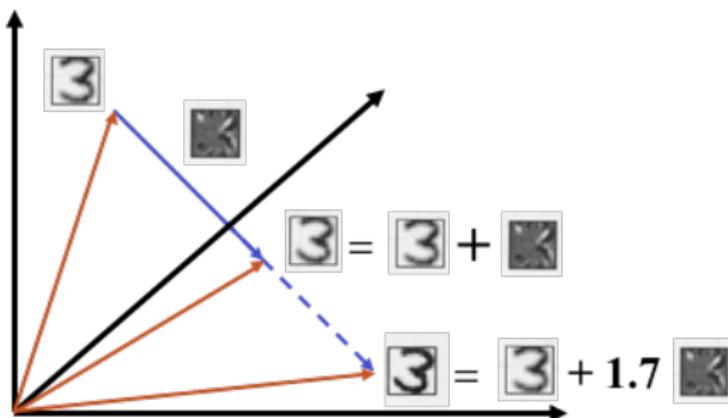
- How can we find more efficient representations for our data?  
How can we **capture the “essence” of the data?**
- Example: images of the digit 3



- The images can be represented as points in a high-dimensional space (e.g., with one dimension per pixel, in a 4k image there are around 9 million dimensions!)

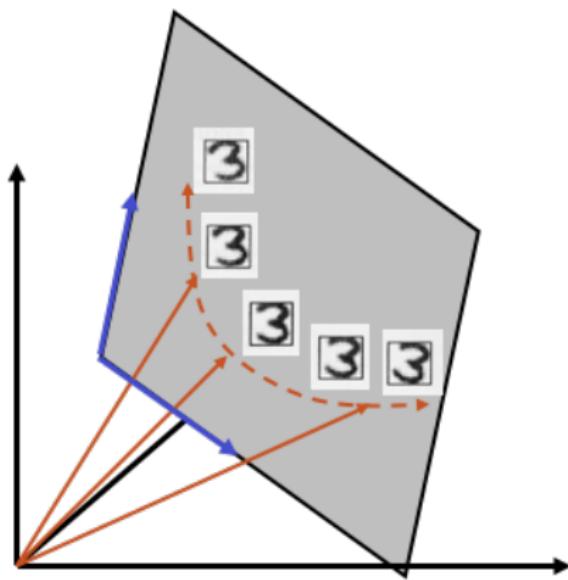
# Linear Dimensionality Reduction

- To make things easier, we will once again assume linear models. A data point (here: one image) can be written as a linear combination of bases (here: basis images)



# Linear Dimensionality Reduction

- What **linear transformations** of the data can be used to define a **lower-dimensional subspace** that captures most of the structure?



# Goal of Linear Dimensionality Reduction

- Original data point  $n$ :  $\mathbf{x}^n \in \mathbb{R}^M$
- Low-dimensional representation of data point  $n$ :  $\mathbf{a}^n \in \mathbb{R}^D$  with  $D \ll M$
- Goal: find a mapping
$$\mathbf{x}^n \rightarrow \mathbf{a}^n$$
- Restrict this mapping to be a linear function

$$\mathbf{a}^n = \mathbf{B}\mathbf{x}^n, \quad \mathbf{B} \in \mathbb{R}^{D \times M}$$

# Simple Observation

- We can always write a vector as

$$\mathbf{x} = \sum_{i=1}^M a_i \mathbf{u}_i, \text{ where } \mathbf{u}_i^\top \mathbf{u}_j = \delta_{ij}$$

- $\delta_{ij} = 1$  if  $i = j$ , 0 otherwise.  $\mathbf{u}_i$  and  $\mathbf{u}_j$  are **orthonormal** to each other

- Example

$$\begin{bmatrix} 3 \\ 7 \end{bmatrix} = 3 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 7 \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

# Simple Observation

$$\begin{aligned}
 \mathbf{x} &= a_1 \mathbf{u}_1 + a_2 \mathbf{u}_2 \\
 a_1 &= \mathbf{u}_1^\top (\mathbf{x} - a_2 \mathbf{u}_2) \\
 &= \mathbf{u}_1^\top \mathbf{x} - \mathbf{u}_1^\top a_2 \mathbf{u}_2 \\
 &= \mathbf{u}_1^\top \mathbf{x} - a_2 \underbrace{\mathbf{u}_1^\top \mathbf{u}_2}_{=0} \\
 &= \mathbf{u}_1^\top \mathbf{x}
 \end{aligned}$$

- More generally

$$\underbrace{a_i}_{\text{scalar coefficient}} = \underbrace{\mathbf{u}_i^\top \mathbf{x}}_{\text{projection}}$$

# Decomposition

$$\mathbf{x}^n = \sum_{i=1}^D a_i \mathbf{u}_i + \underbrace{\sum_{j=D+1}^M b_j \mathbf{u}_j}_{\text{error}} \approx \tilde{\mathbf{x}}^n$$

- We want the  $D$  bases that minimize the mean squared error over the training data

$$\mathbf{u}_1, \dots, \mathbf{u}_D = \arg \min_{\mathbf{u}_1, \dots, \mathbf{u}_D} E(\mathbf{u}_1, \dots, \mathbf{u}_D) = \arg \min_{\mathbf{u}_1, \dots, \mathbf{u}_D} \sum_{n=1}^N \|\mathbf{x}^n - \tilde{\mathbf{x}}^n\|^2$$

# Minimizing the error

- Rewrite the error (assuming a single basis vector)

$$\begin{aligned} E(\mathbf{u}) &= \sum_{n=1}^N \|\mathbf{x}^n - \tilde{\mathbf{x}}^n\|^2 \\ &= \sum_{n=1}^N \|\mathbf{x}^n - (\mathbf{u}^\top \mathbf{x}^n) \mathbf{u}\|^2 \\ &= \sum_{n=1}^N \|\mathbf{x}^n\|^2 - 2(\mathbf{u}^\top \mathbf{x}^n)^2 + (\mathbf{u}^\top \mathbf{x}^n)^2 \cdot \mathbf{u}^\top \mathbf{u} \\ &= \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (\mathbf{u}^\top \mathbf{x}^n)^2 \\ &= \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (a_n)^2 \end{aligned}$$

# Minimizing the error

- Rewrite the error (assuming a single basis vector)

$$E(\mathbf{u}) = \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (\mathbf{u}^\top \mathbf{x}^n)^2 = \sum_{n=1}^N \|\mathbf{x}^n\|^2 - (a_n)^2$$

- Minimizing the error is equivalent to maximizing the variance of the projection. Assuming a zero mean on the data

$$\max \frac{1}{N} \sum_{n=1}^N a_n^2$$

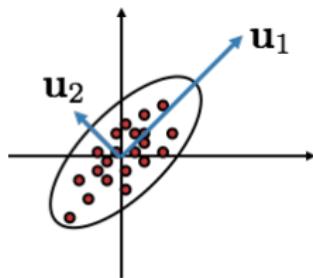
- We can ensure a zero mean projection by subtracting the mean from every data point

$$\mathbf{x}^n - \bar{\mathbf{x}}$$

# Intuition

$$\mathbf{x}^n - \bar{\mathbf{x}} = \sum_{i=1}^D a_i \mathbf{u}_i + \sum_{j=D+1}^M b_j \mathbf{u}_j$$

$$\tilde{\mathbf{x}}^n = \sum_{i=1}^D a_i \mathbf{u}_i + \bar{\mathbf{x}}$$

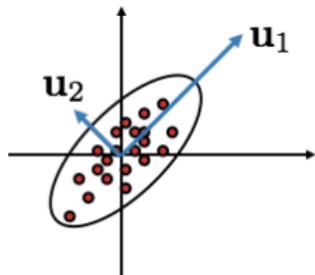


- Projecting onto  $\mathbf{u}_1$  captures the majority of the variance and hence projecting onto it minimizes the error
- What is that and how do we find the axis of largest variance?

# Intuition

$$\mathbf{x}^n - \bar{\mathbf{x}} = \sum_{i=1}^D a_i \mathbf{u}_i + \sum_{j=D+1}^M b_j \mathbf{u}_j$$

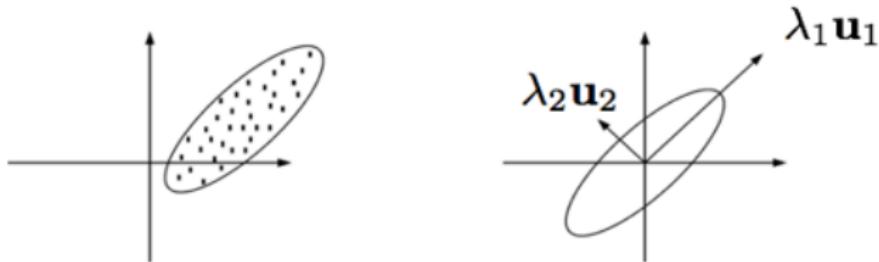
$$\tilde{\mathbf{x}}^n = \sum_{i=1}^D a_i \mathbf{u}_i + \bar{\mathbf{x}}$$



- Note that these axes are **orthogonal** and **decorrelate** the data; i.e. in the coordinate frame of these axes, the **data is uncorrelated** (side note: this only works for Gaussians)

# Principal Component Analysis

- Goal: find the so-called **principal directions**, and the **variance of the data** along each principal direction



- $\lambda_i$  is the **marginal variance** along the **principal direction**  $u_i$

# Principal Component Analysis

- The first principal direction  $u_1$  is the direction along which the variance of the data is maximal

$$u_1 = \arg \max_u u^T \mathbf{C} u \quad \text{s.t. } u^T u = 1$$

- The second principal direction maximizes the variance of the data in the orthogonal complement of the first principal direction

# Problem Formulation

- Let  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^n] \in \mathbb{R}^{M \times N}$  be a matrix of  $N$  vectors in a  $M$ -dimensional input space, with  $\mathbf{x}^i \in \mathbb{R}^M$
- Let  $\mathbf{u} \in \mathbb{R}^M$  be a direction (a vector of length 1) in the input space
- The projection of the  $j$ -th vector  $\mathbf{x}^j$  onto the vector  $\mathbf{u}$  can be computed as

$$a_j = \mathbf{u}^\top \mathbf{x}^j = \sum_{i=1}^M X_{ij} u_i$$

- The goal is to find a direction  $\mathbf{u}$  that maximizes the variance of the projections of all input vectors

$$\mathbf{x}^j \quad \forall j = 1, \dots, N$$

# Variance of the Projection

$$\bar{a} = \frac{1}{N} \sum_{j=1}^N a_j = \frac{1}{N} \sum_{j=1}^N \sum_{i=1}^M \mathbf{x}_{ij} u_i = \sum_{i=1}^M u_i \mu_i$$

where  $\mu_i = \frac{1}{N} \sum_{j=1}^N X_{ij}$  and  $\boldsymbol{\mu} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_j$

$$\begin{aligned}\sigma^2 &= \frac{1}{N} \sum_{j=1}^N (a_j - \bar{a})^2 = \frac{1}{N} \sum_{j=1}^N \left( \sum_{i=1}^M u_i \mathbf{x}_{ij} - \sum_{i=1}^M u_i \mu_i \right)^2 \\ &= \frac{1}{N} \sum_{j=1}^N \left( \sum_{i=1}^M u_i \hat{\mathbf{x}}_{ij} \right)^2 = \frac{1}{N} \sum_{j=1}^N \sum_{i=1}^M \sum_{k=1}^M u_i \hat{\mathbf{x}}_{ij} u_k \hat{\mathbf{x}}_{kj} \\ &= \sum_{i=1}^M \sum_{k=1}^M u_i u_k \underbrace{\frac{1}{N} \langle \hat{\mathbf{x}}_{i:}, \hat{\mathbf{x}}_{k:} \rangle}_{\mathbf{C} = \frac{1}{N} \hat{\mathbf{x}} \hat{\mathbf{x}}^\top} = \sum_{i=1}^M \sum_{k=1}^M u_i C_{ik} u_k = \mathbf{u}^\top \mathbf{C} \mathbf{u}\end{aligned}$$

- $\mathbf{C}$  is the covariance matrix

# Maximizing the Variance

- Maximize  $\sigma^2$  under the constraint that  $\|\mathbf{u}\| = 1$
- We can translate into an optimization with Lagrangian multipliers

$$F(\mathbf{u}; \lambda) = \mathbf{u}^\top \mathbf{C} \mathbf{u} - \lambda \left( \sum_{k=1}^M u_k^2 - 1 \right) = \sum_{i=1}^M \sum_{j=1}^M u_i C_{ij} u_j - \lambda \left( \sum_{k=1}^M u_k^2 - 1 \right)$$

$$\frac{\partial F}{\partial u_l} = \sum_{j=1}^M C_{lj} u_j + \sum_{i=1}^M u_i C_{il} - \lambda 2u_l = 0 \quad \forall l = 1, \dots, M$$

$$\sum_{j=1}^M C_{lj} u_j = \lambda u_l \quad \forall l = 1, \dots, M \iff \mathbf{C} \mathbf{u} = \lambda \mathbf{u}$$

# Maximizing the Variance

$$\sum_{j=1}^M C_{lj} u_j = \lambda u_l \quad \forall l = 1, \dots, M \iff \mathbf{C}\mathbf{u} = \lambda\mathbf{u}$$

- Equation of Eigenvalues-Eigenvectors
- The largest Eigenvalue gives us the maximal variance
- The corresponding Eigenvector gives us the direction with maximal variance

# Eigendecomposition - Reminder

- Let the Eigenvectors and Eigenvalues of  $\mathbf{C}$  be  $\mathbf{u}_k$  and  $\lambda_k$  for  $k \leq M$ , i.e.,  $\mathbf{C}\mathbf{u}_k = \lambda_k\mathbf{u}_k$  with  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_M$
- In matrix form:  $\mathbf{C}\mathbf{U} = \mathbf{U}\Lambda$ , where  $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_M)$  and  $\mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_M]$
- Because  $\mathbf{U}$  is orthonormal (we assume that the Eigenvectors have unit norm) we know that  $\mathbf{U}\mathbf{U}^\top = \mathbf{I}$
- This means that we can decompose  $\mathbf{C}$  as

$$\mathbf{C} = \mathbf{U}\Lambda\mathbf{U}^{-1}$$

$$\mathbf{C} = \mathbf{U}\Lambda\mathbf{U}^\top$$

# Principal Component Analysis

- By definition of a covariance matrix,  $\mathbf{C}$  is **real, symmetric** and **positive-definite**. Thus we can decompose it in its **Eigendecomposition** as

$$\mathbf{C} = \mathbf{U}\Lambda\mathbf{U}^T = \underbrace{\begin{bmatrix} \mathbf{u}_1 & \mathbf{u}_2 & \dots & \mathbf{u}_M \end{bmatrix}}_{\text{Eigenvectors}} \underbrace{\begin{bmatrix} \lambda_1 & & & \\ & \lambda_2 & & \\ & & \ddots & \\ & & & \lambda_M \end{bmatrix}}_{\text{Eigenvalues}} \underbrace{\begin{bmatrix} \mathbf{u}_1^T \\ \mathbf{u}_2^T \\ \vdots \\ \mathbf{u}_M^T \end{bmatrix}}_{\mathbf{U}^T}$$

# Principal Component Analysis

- **Observation:** If  $\lambda_k \approx 0$  for  $k > D$  for some  $D \ll M$ , then we can use the subset of the first  $D$  eigenvectors to define a basis for approximating the data vectors

$$\mathbf{x}^n - \bar{\mathbf{x}} = \sum_{i=1}^D a_i \mathbf{u}_i + \sum_{j=D+1}^M b_j \mathbf{u}_j$$

$$\mathbf{x}^n \approx \tilde{\mathbf{x}}^n = \bar{\mathbf{x}} + \sum_{i=1}^D a_i \mathbf{u}_i \quad \text{where } a_i = \mathbf{u}_i^\top (\mathbf{x}^n - \bar{\mathbf{x}})$$

- This representation has the **minimal mean squared error** (MSE) of all linear representations of dimension  $D$

$$\min E(\mathbf{u}_1, \dots, \mathbf{u}_D) = \sum_{n=1}^N \|\mathbf{x}^n - \tilde{\mathbf{x}}^n\|^2$$

# Principal Component Analysis

- Now we know how we can represent our data in a lower dimensional space in a principled way
  - Center the data around the mean (compute the mean of the data and subtract it)
  - Compute the covariance matrix, decompose it, and choose the first  $D$  largest Eigenvalues and corresponding Eigenvectors
  - This gives us an (Eigen)basis for representing the data

$$\begin{aligned}\mathbf{a}^n &= \mathbf{B}^T (\mathbf{x}^n - \bar{\mathbf{x}}) \\ \tilde{\mathbf{x}}^n &= \bar{\mathbf{x}} + \mathbf{B}\mathbf{a}^n\end{aligned}$$

where  $\mathbf{B} = [\mathbf{u}_1 \dots \mathbf{u}_D]$

- It is also common to normalize the variance of each dimension
- But how to choose  $D$ ?

# Choosing $D$

- A larger  $D$  leads to a better approximation. In the limit, when  $D = M$  we stay in the initial data dimensions
- There are at least 2 good possibilities for choosing  $D$

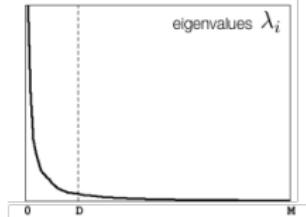
# Choosing $D$

- A larger  $D$  leads to a better approximation. In the limit, when  $D = M$  we stay in the initial data dimensions
- There are at least 2 good possibilities for choosing  $D$ 
  1. Choose  $D$  based on application performance, i.e. **choose the smallest  $D$  that makes the application work well enough**

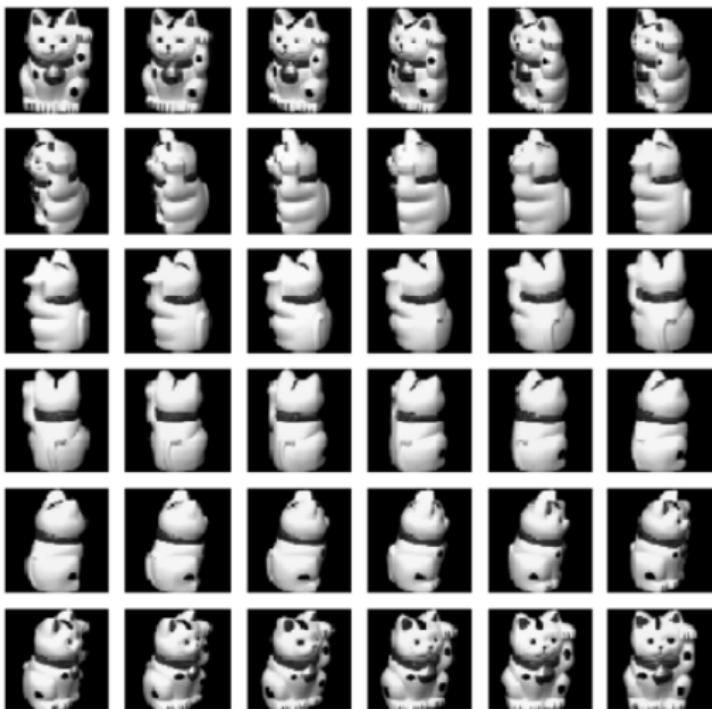
# Choosing $D$

- A larger  $D$  leads to a better approximation. In the limit, when  $D = M$  we stay in the initial data dimensions
- There are at least 2 good possibilities for choosing  $D$ 
  1. Choose  $D$  based on application performance, i.e. choose the smallest  $D$  that makes the application work well enough
  2. Choose  $D$  so that the eigenbasis captures some fraction of the variance (for example  $\eta = 0.9$ )  
The eigenvalue  $\lambda_i$  describes the marginal variance captured by  $\mathbf{u}_i$

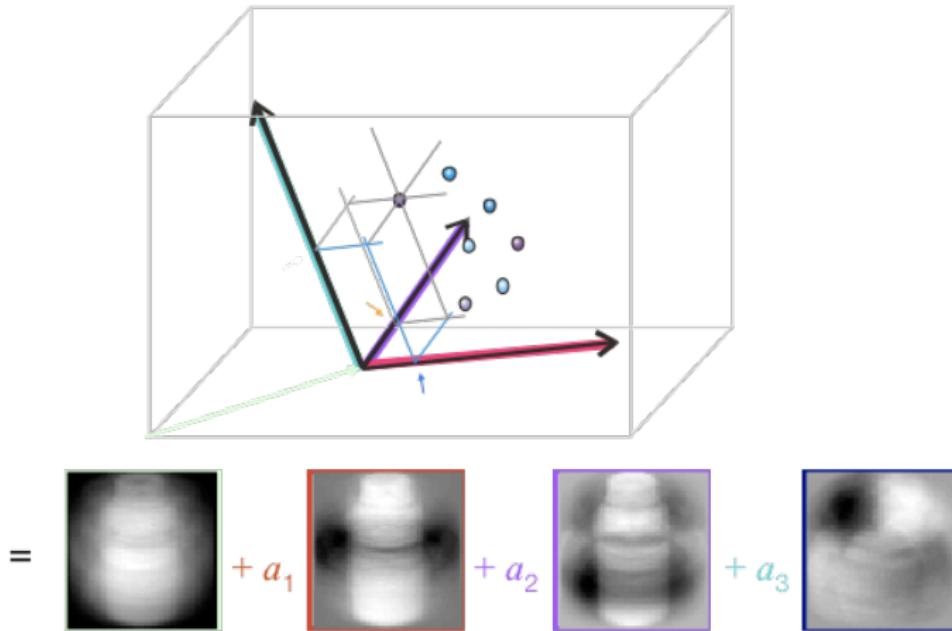
Choose  $D$  s.t.  $\sum_{i=1}^D \lambda_i \geq \eta \sum_{i=1}^M \lambda_i$



# Image Representation

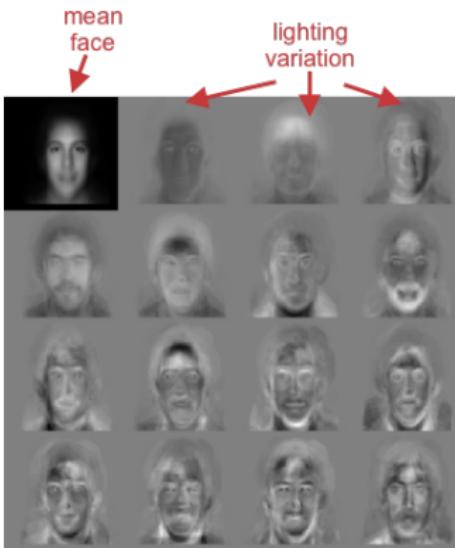


# Image Representation with PCA

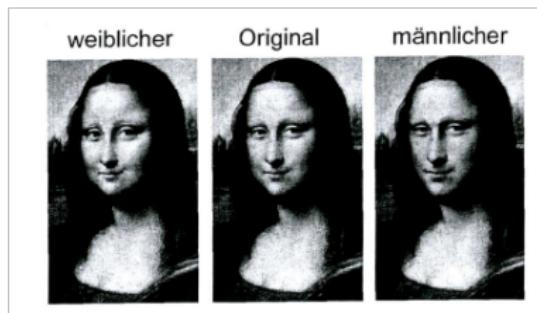


# EigenFaces

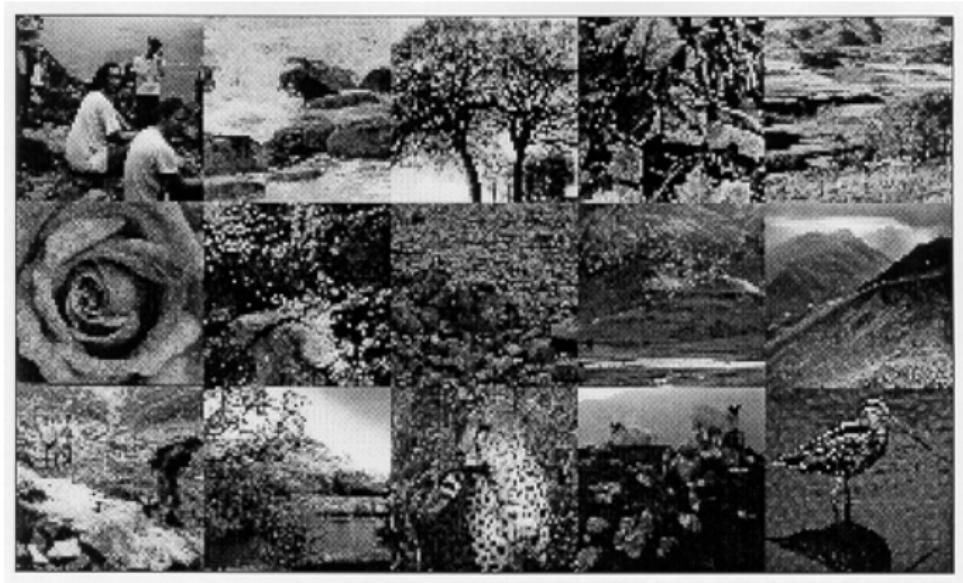
- The first popular use of PCA for object recognition was for the detection and recognition of faces [Turk and Pentland, 1991]
- Collect a face ensemble
- Normalize for contrast, scale, & orientation
- Remove backgrounds
- Apply PCA & choose the first  $D$  eigen-images that account for most of the variance of the data



# Morphing with EigenFaces



# Generic Image Ensembles



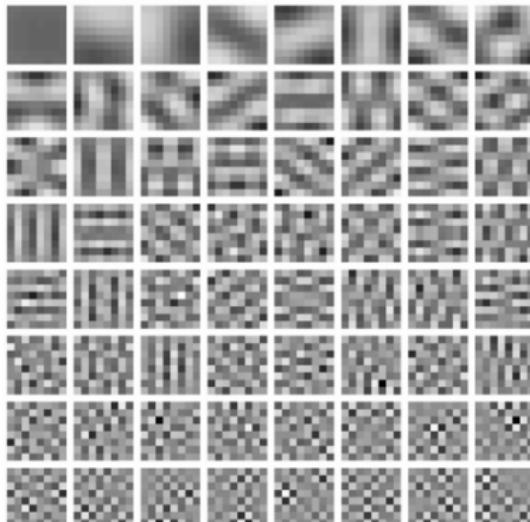
# Generic Image Ensembles



Is there a low dimensional model describing natural images?

# PCA of Natural Image Patches

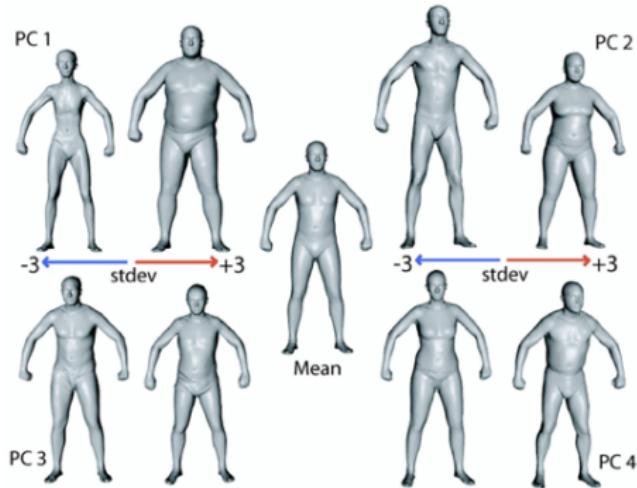
8x8 image patches



What do these bases look like?

# PCA Model of Body Shapes

PCA on a detailed triangle model of human bodies [Anguelov et al. 05]



# PCA Summary and Applications

## ■ Summary

- PCA projects the data into a linear subspace
- PCA maximizes the variance of the projection
- PCA minimizes the error of the reconstruction

## ■ Applications

- PCA allows us to transform a high-dimensional input space to a low-dimensional feature space, while capturing the essence of the data
- PCA finds a more *natural* coordinate system for the data
- PCA is a very common **preprocessing step for high-dimensional input** data

# Outline

1. Linear Dimensionality Reduction

**2. Statistical Learning Theory**

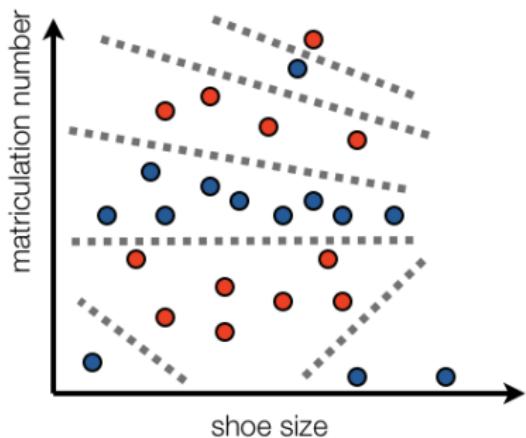
3. Wrap-Up

# Our approach to Machine Learning until now

- We have mostly dealt with *Classical Statistical Learning*
  - Estimate a parameter vector  $\mathbf{w}$  of a fixed model (*learning machine*)
  - Examples
    - Maximum Likelihood estimation: determine the parameters that model a density function
    - Discriminant functions: determine the orientation of the separating hyperplane
- Learning occurs by optimizing the model parameters
- In this setting we assume that we already know in advance the correct model (Bayesian learning is the exception)
  - For instance we model the density function with a linear combination of features, a neural network, a Gaussian, ...

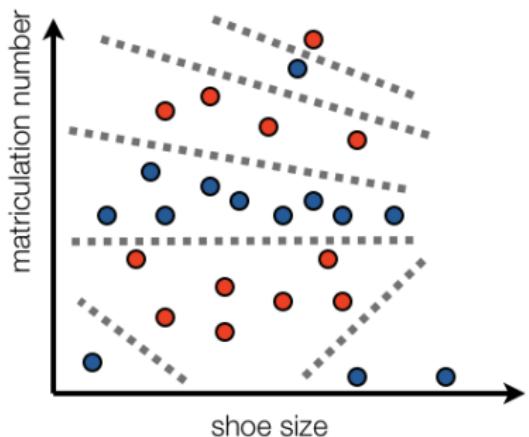
# Generalization Abilities

- Example: Classify students who will pass the exam. What are some common challenges?



# Generalization Abilities

- Example: Classify students who will pass the exam. What are some common challenges?



- Problem 1: Possibly irrelevant features
- Problem 2: Overfitting from overly complex model
- We are really interested in the **generalization ability** and the corresponding risk

# Statistical Learning Theory

- Statistical Learning Theory
  - Does **not assume that the correct model is known in advance**
  - The goal is to choose an *optimal model* from a set of specified models
  - It is a form of *model selection*
- Optimality here means
  - Ability of the model to generalize, i.e. to have the **lowest error probability on all data**, and not only on the training data

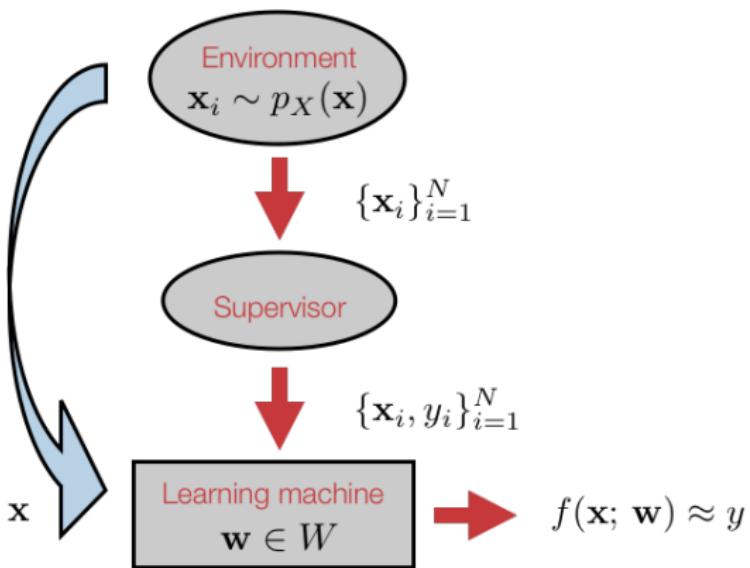
# Statistical Learning Theory



Vladimir Vapnik [1936-]

- Statistical Learning Theory (from Vladimir Vapnik)
  - Is concerned with the question how one can control the generalization abilities of a learning machine
  - Aims at a **formal description** of the generalization ability
  - The goal is to develop a **rigorous theory** as opposed to commonly used heuristics
- We will see later that this may be a noble goal, but the theory unfortunately does not say that much about real problems...

# Supervised Learning



# Statistical Learning Theory

- Supervised learning as an example
  - The environment is stationary, i.e., the data points  $\mathbf{x}_i$  have an unknown, but fixed probability density
$$\mathbf{x}_i \sim p_{\mathbf{X}}$$
  - The supervisor returns the intended classification label  $y$  for every data point  $\mathbf{x}$ , possibly with some added noise  $v$ 
$$y = g(\mathbf{x}, v)$$
  - The learning machine is represented through a class of functions (with parameters  $\mathbf{w}$ ) that return an output  $y$  for every input  $\mathbf{x}$ 
$$y = f(\mathbf{x}, \mathbf{w})$$

# Statistical Learning Theory

- Supervised learning from the learning machine view
  - Choose a particular function

$$y = f(\mathbf{x}, \mathbf{w})$$

- Given a set of training examples

$$\{\mathbf{x}_i, y_i\}_{i=1}^N$$

- Goal: the desired output  $y$  should be approximated *optimally*
- Assessment of *optimality*
  - A loss function, for example the quadratic loss

$$L(y, f(\mathbf{x}, \mathbf{w})) = (y - f(\mathbf{x}, \mathbf{w}))^2$$

# Assessment of optimality - Risk

- Loss function

$$L(y, f(\mathbf{x}, \mathbf{w}))$$

- Empirical risk

$$R_{\text{emp}}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(\mathbf{x}_i, \mathbf{w}))$$

where  $N$  is the number of data points

- Example with the quadratic loss function

$$R_{\text{emp}}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i, \mathbf{w}))^2$$

# Risk

- In reality, we are instead interested in the **True Risk**

$$\begin{aligned} R(\mathbf{w}) &= \int L(y, f(\mathbf{x}, \mathbf{w})) p(\mathbf{x}, y) d\mathbf{x}dy \\ &= \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [L(y, f(\mathbf{x}, \mathbf{w}))] \end{aligned}$$

where  $p(\mathbf{x}, y)$  is the **true joint probability density** of  $\mathbf{x}$  and  $y$

- The risk is the expected error over all data sets
- The risk is the expectation of the generalization error
- What is the problem here?

# Risk

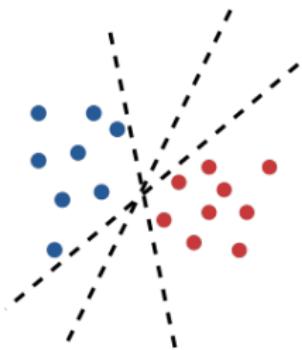
- In reality, we are instead interested in the **True Risk**

$$\begin{aligned} R(\mathbf{w}) &= \int L(y, f(\mathbf{x}, \mathbf{w})) p(\mathbf{x}, y) d\mathbf{x}dy \\ &= \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [L(y, f(\mathbf{x}, \mathbf{w}))] \end{aligned}$$

where  $p(\mathbf{x}, y)$  is the **true joint probability density** of  $\mathbf{x}$  and  $y$

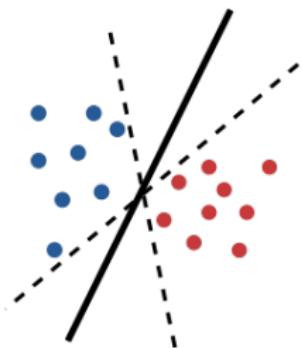
- The risk is the expected error over all data sets
- The risk is the expectation of the generalization error
- What is the problem here?
  - $p(\mathbf{x}, y)$  is fixed but usually unknown
  - We cannot compute the (actual) risk directly

# Empirical vs True Risk



- All 3 decision boundaries have zero empirical risk
- Which one is preferable?
- Which one generalizes best?

# Empirical vs True Risk



- All 3 decision boundaries have zero empirical risk
- Which one is preferable?
- Which one generalizes best?

# Empirical vs True Risk

$$R(\mathbf{w}) = \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [L(y, f(\mathbf{x}, \mathbf{w}))], \quad R_{\text{emp}}(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N L(y_i, f(\mathbf{x}_i, \mathbf{w}))$$

## ■ True risk

- Advantage: Measure for the generalization ability
- Disadvantage: In general, we do not know  $p(\mathbf{x}, y)$

## ■ Empirical risk

- Disadvantage: No direct measure for the generalization ability
- Advantage: Does not depend on  $p(\mathbf{x}, y)$
- Learning algorithms often minimize the empirical risk

- We are interested in the dependencies between these two risks

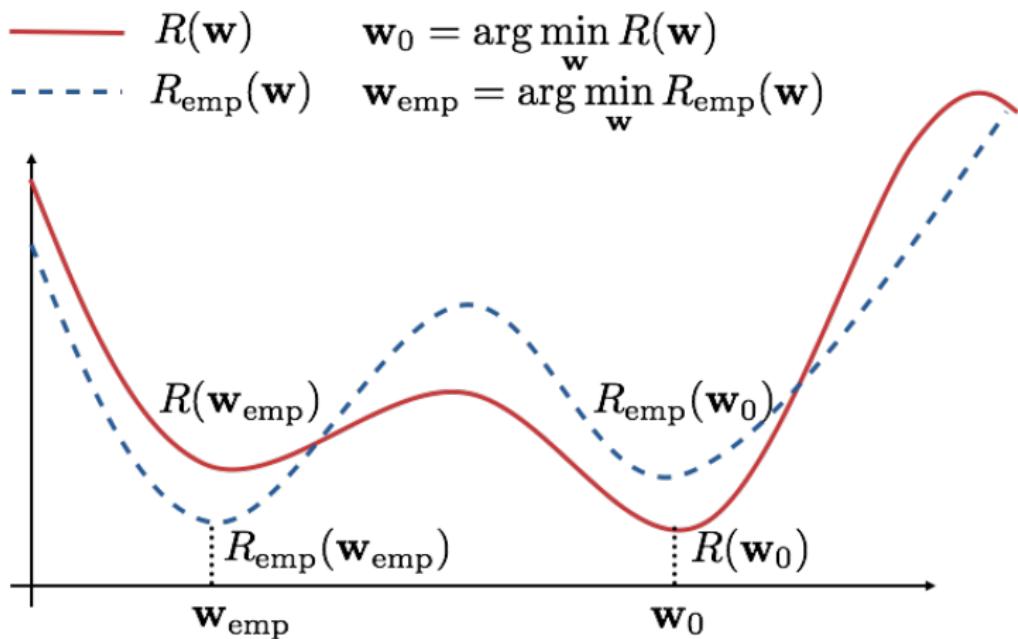
# Empirical vs True Risk

- You may be asking yourself: we learned that with Monte Carlo we can estimate an expectation by sampling and averaging. So, isn't the empirical risk the same as the true risk?

$$R(\mathbf{w}) = \mathbb{E}_{\mathbf{x}, y \sim p(\mathbf{x}, y)} [L(y, f(\mathbf{x}, \mathbf{w}))] \approx \frac{1}{N} \sum_{i=1}^N L(y_i, f(\mathbf{x}_i, \mathbf{w})) = R_{\text{emp}}(\mathbf{w})$$

- It is just an approximation, which works very well if our distribution is very concentrated
- The approximation is very good if we have infinitely many samples from the joint distribution
- We will see now how the two risks relate

# Convergence Properties



# Convergence Properties - Derivation

- We assume that the empirical risk converges to the true risk as we supply more and more training data

$$\inf_{\mathbf{w}} R_{\text{emp}}(\mathbf{w}) \xrightarrow{P} \inf_{\mathbf{w}} R(\mathbf{w}) \quad \text{as } N \rightarrow \infty$$

where  $\inf$  stands for *infimum*

# Convergence Properties - Derivation

- We assume that the empirical risk converges to the true risk as we supply more and more training data

$$\inf_{\mathbf{w}} R_{\text{emp}}(\mathbf{w}) \xrightarrow{P} \inf_{\mathbf{w}} R(\mathbf{w}) \quad \text{as } N \rightarrow \infty$$

where  $\inf$  stands for *infimum*

- We further assume that the convergence has to be uniform

$$P \left( \sup_{\mathbf{w}} |R(\mathbf{w}) - R_{\text{emp}}(\mathbf{w})| > \epsilon \right) = 0 \quad \text{as } N \rightarrow \infty$$

where  $\sup$  stands for *supremum*

- Intuitively, “the learning machine gets better as we give it more training data”

# Convergence Properties - Derivation

- If the convergence is uniform

$$P \left( \sup_{\mathbf{w}} |R(\mathbf{w}) - R_{\text{emp}}(\mathbf{w})| > \epsilon \right) < p^* \quad \text{for some } p^* > 0$$

- Then with probability  $(1 - p^*)$  it holds that

$$\begin{aligned} |R(\mathbf{w}_{\text{emp}}) - R_{\text{emp}}(\mathbf{w}_{\text{emp}})| &< \epsilon \\ |R(\mathbf{w}_0) - R_{\text{emp}}(\mathbf{w}_0)| &< \epsilon \end{aligned}$$

- Hence we have

$$P(|R(\mathbf{w}_0) - R_{\text{emp}}(\mathbf{w}_{\text{emp}})| > 2\epsilon) < p^*$$

# Empirical Risk Minimizing Principle

- Necessary and sufficient condition: uniform convergence

$$P \left( \sup_{\mathbf{w}} |R(\mathbf{w}) - R_{\text{emp}}(\mathbf{w})| > \epsilon \right) = 0 \quad \text{as } N \rightarrow \infty$$

- Compute the empirical risk from the available training data  
 $\{\mathbf{x}_i, y_i\}_{i=1}^N$

$$R_{\text{emp}} = \frac{1}{N} \sum_{i=1}^N L(y_i, f(\mathbf{x}_i, \mathbf{w}))$$

- Minimizing the empirical risk guarantees that we minimize the true risk in the limit of  $N \rightarrow \infty$

# Empirical Risk Minimizing Principle

## ■ Advantages

- We have a **formal criterion** as to what we can expect in terms of generalization
- Necessary and sufficient condition is the **uniform convergence**

## ■ Disadvantages

- In reality, the **amount of training data is very limited**
- We can't easily "take the limit" of  $N \rightarrow \infty$
- Nonetheless, what we have essentially been assuming so far in many cases

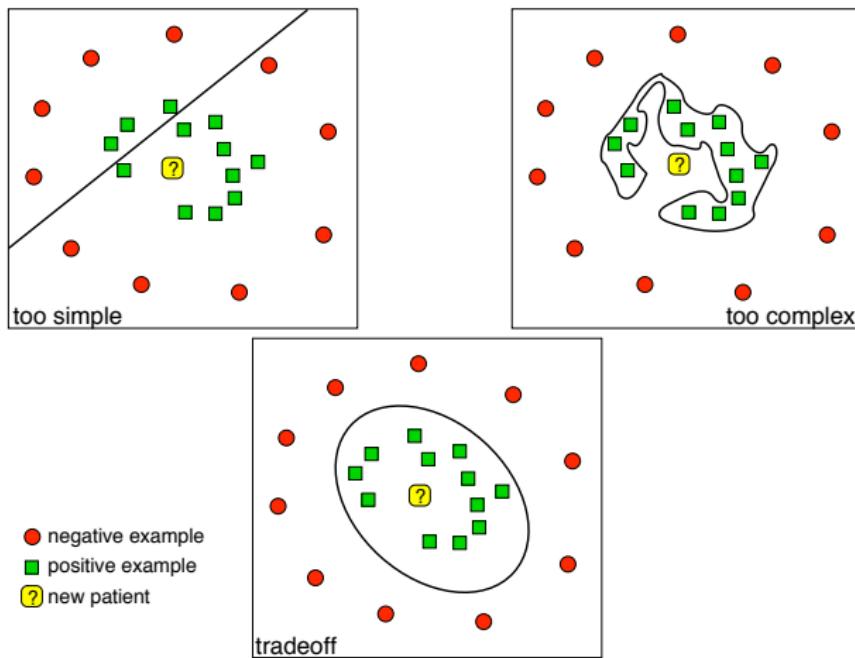
# Risk Bound

- Idea: determine an upper bound on the true risk based on the empirical risk

$$R(\mathbf{w}) \leq R_{\text{emp}}(\mathbf{w}) + \epsilon(N, p^*, h)$$

- $N$  is the number of training examples
- $p^*$  is the probability that the bound is met
- $h$  is the *learning power* of the learning machine, formally called the **VC-dimension**

# Learning Power



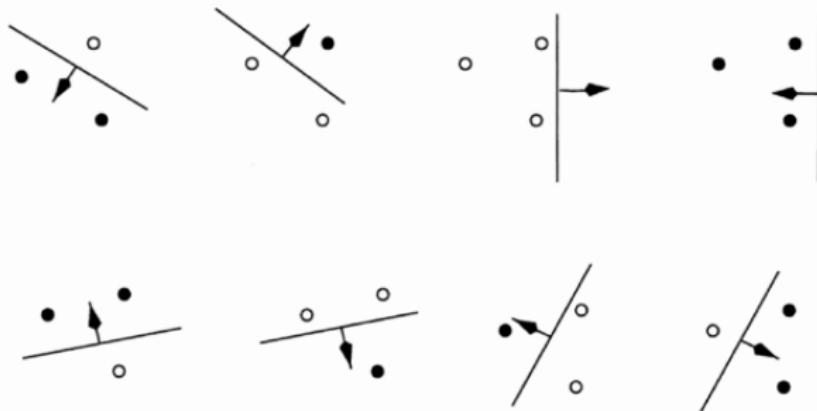
[Florian Markowetz]

# VC-Dimension

- VC stands for Vapnik-Chervonenkis
- Informal definition of the VC-dimension
  - The VC-dimension of a family of functions is the maximal number of data points that can be correctly classified by a function from that family (no matter which label configuration the data points have)
  - The VC-dimension is a measure for the capacity, or “learning power”, of a classifier

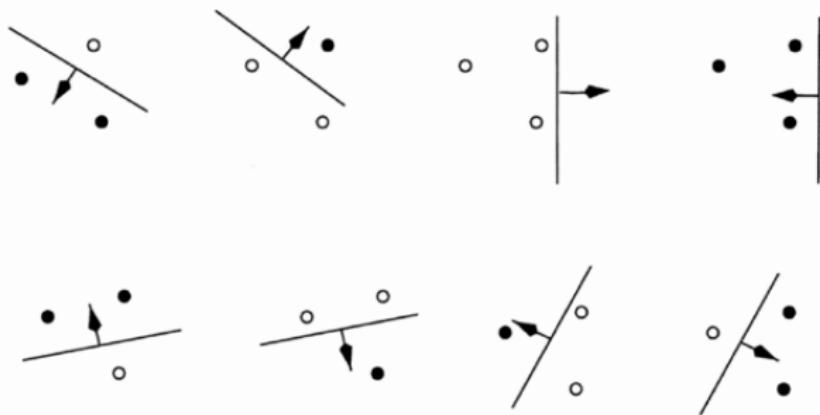
# VC-Dimension Example

- The VC-dimension of linear classifiers (lines) in  $\mathbb{R}^2$  is 3



# VC-Dimension Example

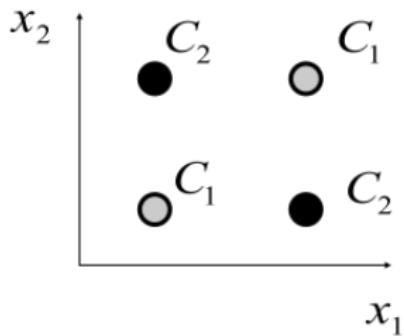
- The VC-dimension of linear classifiers (lines) in  $\mathbb{R}^2$  is 3



- More generally, the VC-dimension of linear classifiers (hyperplanes) in  $\mathbb{R}^n$  is  $n + 1$ .

# VC-Dimension Example

- The VC-dimension of linear classifiers (lines) in  $\mathbb{R}^2$  is 3. Why is it not bigger than 3?



# VC-Dimension

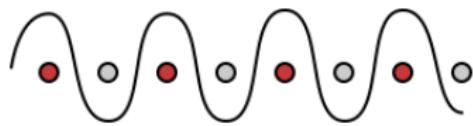
- Intuition (but unfortunately not quite right)
  - The VC-dimension is directly related to the number of parameters
- Counter example

$$\begin{aligned}f(x, \mathbf{w}) &= g(\sin(w_1 x + w_0)) \\g(x) &= \begin{cases} 1 & x > 0 \\ -1 & x \leq 0 \end{cases}\end{aligned}$$

- Has only two parameters but infinite VC-dimension

# VC-Dimension Counter Example

$$\begin{aligned}f(x, \mathbf{w}) &= g(\sin(w_1 x + w_0)) \\g(x) &= \begin{cases} 1 & x > 0 \\ -1 & x \leq 0 \end{cases}\end{aligned}$$



- For other patterns of data, there is always a setting that correctly classifies any labeling of a fixed number of points

# Risk Bound Example

- Loss function

$$L(y, f(\mathbf{x}, \mathbf{w})) = \frac{1}{2} |y - f(\mathbf{x}, \mathbf{w})|$$

- True risk

$$R(\mathbf{w}) = \int \frac{1}{2} |y - f(\mathbf{x}, \mathbf{w})| p(\mathbf{x}, y) d\mathbf{x} dy$$

- Empirical risk

$$R_{\text{emp}}(\mathbf{w}) = \frac{1}{2N} \sum_{i=1}^N |y_i - f(\mathbf{x}_i, \mathbf{w})|$$

# Risk Bound Example

- With probability  $p^*$  it holds that

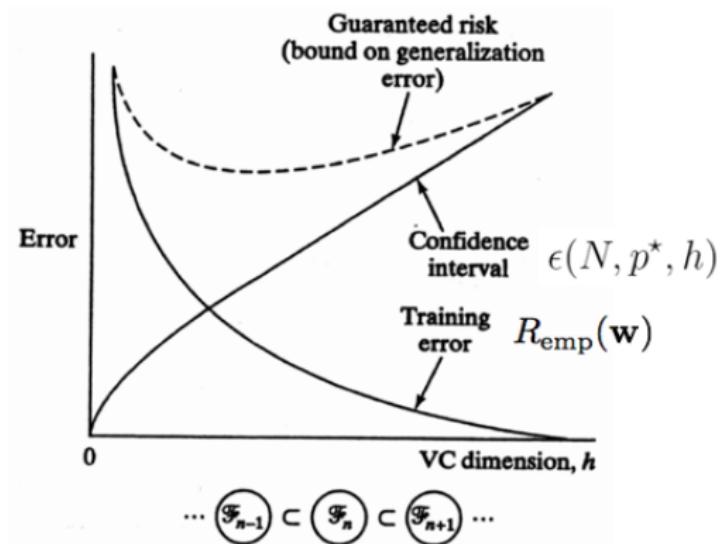
$$R(\mathbf{w}) \leq R_{\text{emp}}(\mathbf{w}) + \sqrt{\frac{h(\log(2N/h) + 1) - \log((1-p^*)/4)}{N}}$$

- The upper bound is independent of  $p(\mathbf{x}, y)$
- Typically we cannot compute the true risk
- But if we know the VC-dimension, we can compute a bound of the type

$$R(\mathbf{w}) \leq R_{\text{emp}}(\mathbf{w}) + \epsilon(N, p^*, h)$$

- Note, however, that in practice this bound is very loose

# Risk Bound



# Structural Risk Minimization

- Given a family of  $n$  models  $f_i(\mathbf{x}_i; \mathbf{w}_i)$  with non-decreasing VC-dimension

$$h_1 \leq h_2 \leq \dots \leq h_n$$

- Minimize the empirical risk for every model
- Choose the model that minimizes the risk bound (Right Hand Side of the true/empirical risk inequality)
- In general, this is not the same model that minimizes the empirical risk
- This is formally justified by the bound on the true risk
- The result is only sensible, however, if the upper bound on the true risk is a tight bound

# Outline

1. Linear Dimensionality Reduction

2. Statistical Learning Theory

**3. Wrap-Up**

## 3. Wrap-Up

You know now:

- What dimensionality reduction is and why do we need it
- What the intuition behind PCA is
- Why maximizing the variance of the projection is a good idea
- How PCA relates to Eigenvectors and Eigenvalues
- What Statistical Learning Theory is
- What empirical and true risk are and how they relate
- Why looking just at the empirical risk is not enough
- What the VC-Dimension is
- How to relate the VC-Dimension with the model complexity

# Self-Test Questions

- What does dimensionality reduction mean?
- What is PCA? What are the three things that it does?
- What are the roles of the Eigenvectors and Eigenvalues in PCA?
- Can you describe applications of PCA?
- What does risk in statistical learning theory mean?
- How is the true risk different from the empirical risk?
- What is the learning power of a function approximator?
- What is expressed by a VC-Dimension?
- Is the VC-Dimension always correlated with the number of parameters?

# Homework

- Reading Assignment for next lecture
  - Bishop 5.1, 5.2, 5.3